



Introduction to Pandas

Tabular data and more

Programme: four blocks

Introduction to Series
and DataFrames

First operations,
reading/writing

Plotting capacities

Timeseries analysis



```
| import pandas as pd
```

Part I: pandas objects

Series and Dataframes

pd.Series: introduction

- List of values with indices

Indices of the values

```
Elf  
Dwarf  
Hobbit
```

List of values for a given variable

```
200  
120  
110
```

```
Elf      220  
Dwarf    120  
Hobbit   110  
dtype: int64
```

pd.Series: index and data types

- Variables can have a **type** or not

index 0	4
index 1	7
index 2	5
index 3	7
index 4	2
dtype: int32	

index 0	1.684448
index 1	0.592174
index 2	0.938397
index 3	4.136875
index 4	1.495691
dtype: float64	

index 0 These
index 1 are
index 2 different
index 3 strings
index 4 !
dtype: string

index 0	Hello !
index 1	0.41
index 2	12
index 3	[1, 2, 5]
index 4	x ²
dtype: object	

- Indices can also have different **types**

index 0	3
index 1	6
index 2	2
index 3	1
index 4	0
dtype: int32	

0	8
1	9
2	1
3	1
4	5
dtype: int32	

2012-01-01	9
2013-01-01	5
2014-01-01	8
2015-01-01	7
2016-01-01	4
Freq: YS-JAN, dtype: int32	

pd.Series: creating an object

- Required: a list of values + a list of indices
- Use **pd.Series()**

```
characters = ['Elf', 'Dwarf', 'Hobbit']
size = [200, 120, 110]
pd.Series(data=size, index=characters, name='size', dtype='int')
```

Elf	200
Dwarf	120
Hobbit	110
Name:	size, dtype: int32

Required If not specified: Not required
 0,1,2 ,...

pd.Series: access data

Elf	200
Dwarf	120
Hobbit	110
Name: size, dtype: int32	

`size_series.loc['Elf']`

200

`size_series.iloc[1]`

120

- Access data using **indices**

- Access data using **position**

- Access the data as a numpy array

`size_series.values`

`array([200, 120, 110])`

pd.DataFrame: introduction

- Multiple pd.Series **aligned** on the same indices
- Each Series has a name

pd.Series

Index

	size	weight	life_expectancy	home
Elf	200	80	1000	Rivendale
Dwarf	120	120	300	Moria
Hobbit	110	40	120	The Shire

pd.DataFrame: creating an object

Input data:

```
characters = ['Elf', 'Dwarf', 'Hobbit']

size = [200, 120, 110]
weight = [80, 120, 40]
life_expectancy = [1000, 300, 120]
home = ['Rivendale', 'Moria', 'The Shire']
```

Object creation:

```
pd.DataFrame(data={"size":size,
                   "weight":weight,
                   "life_expectancy":life_expectancy,
                   "home":home},
              index=characters)
```

	size	weight	life_expectancy	home
Elf	200	80	1000	Rivendale
Dwarf	120	120	300	Moria
Hobbit	110	40	120	The Shire

pd.DataFrame: creating an object, option 2

Input data:

```
characters = ['Elf', 'Dwarf', 'Hobbit']

data=[[200, 80, 1000, 'Rivendale'],
      [120, 120, 300, 'Moria'],
      [110, 40, 120, 'The Shire']],
```

Object creation:

```
pd.DataFrame(data=[[200, 80, 1000, 'Rivendale'],
                   [120, 120, 300, 'Moria'],
                   [110, 40, 120, 'The Shire']],
             columns=['size','weight','life_expectancy','home'],
             index=characters)
```

	size	weight	life_expectancy	home
Elf	200	80	1000	Rivendale
Dwarf	120	120	300	Moria
Hobbit	110	40	120	The Shire

pd.DataFrame: access the data

- Select data from index/columns or positions

- Using positions:

```
df.iloc[1,3]
```

```
'Moria'
```

- From index/column

```
df.loc['Dwarf','home']
```

```
'Moria'
```

- Multiple selection

```
df.loc[['Dwarf', 'Elf'],['home', 'weight']]
```

home	weight
------	--------

Dwarf	Moria	120
-------	-------	-----

Elf	Rivendale	80
-----	-----------	----

	size	weight	life_expectancy	home
Elf	200	80	1000	Rivendale
Dwarf	120	120	300	Moria
Hobbit	110	40	120	The Shire

pd.DataFrame: access the data

- Access data at one index
 - Returns a pd.Series

```
df.loc['Elf']
```

```
size          200
weight         80
life_expectancy    1000
home        Rivendale
Name: Elf, dtype: object
```

```
df.iloc[1]
```

```
size          120
weight         120
life_expectancy    300
home        Moria
Name: Dwarf, dtype: object
```

- Access data from one pd.Series
 - Returns a pd.Series

```
df.weight
```

```
Elf      80
Dwarf   120
Hobbit   40
Name: weight, dtype: int64
```

```
df['weight']
```

```
df.iloc[:,3]
```

```
Elf        Rivendale
Dwarf      Moria
Hobbit    The Shire
Name: home, dtype: object
```

pd.DataFrame: append data

- Add a column - pd.Series
 - Add the Series as if you want to access the data of a column

```
df['example_name'] = ['Legolas', 'Gimli', 'Frodo']  
df
```

	size	weight	life_expectancy	home	example_name
Elf	200	80	1000	Rivendale	Legolas
Dwarf	120	120	300	Moria	Gimli
Hobbit	110	40	120	The Shire	Frodo

- The values of the series can also be updated this way.

pd.DataFrame: append data

- Add a row
 - Add the row as if you want to access the data of a line

```
df.loc['Human'] = [170, 70, 75, 'Gondor', 'Aragorn']  
df
```

	size	weight	life_expectancy	home	example_name
Elf	200	80	1000	Rivendale	Legolas
Dwarf	120	120	300	Moria	Gimli
Hobbit	110	40	120	The Shire	Frodo
Human	170	70	75	Gondor	Aragorn

- The values of a line can also be updated this way.

Reading/writing files: supported formats

- Pandas supports many **tabular-like** data formats
 - csv
 - excel (xls, xlsx) : *requiert openpyxl*
 - txt
 - json
 - ...
- Everything that looks like tabular data
- Can open directly from internet using url

Reading/writing files: reading

- To **open** a file, use `pd.read_format()`

```
penguins = pd.read_json('../data/penguins_dataset.json')
planets = pd.read_excel('../data/exoplanets_discoveries.xlsx')
boats = pd.read_csv('../data/fishing_boats.csv')

| pd.read_csv('https://odre.opendatasoft.com/temperature-quotidienne-regionale.csv')
```

- Some important **arguments**

- `index_col` → columns corresponding to index
- `skiprows` → skip the first n rows
- `nrows` → only read the first n rows avec the skipped rows

```
boats = pd.read_csv('../data/fishing_boats.csv', index_col=1, skiprows=5, nrows=10)
```

Reading/writing files: writing

- To **write** a file, use `pd.to_format()`

```
personnages_lotr.to_csv('..../data/lotr_personnages.csv')
```

```
personnages_lotr.to_json('..../data/lotr_personnages.json')
```

```
personnages_lotr.to_excel('..../data/lotr_personnages.xlsx')
```

Part I: Summary

pd.Series : one variable only

pd.DataFrame : multiple variables aligned

- Easy access to data using **index** values
- Objects are flexible and can be **modified**

Part I: Practicals

Go to the jupyter notebook



Part II: First analyses

Statistical operations, filtering, groups and others

Filter data

Data where mean radius > 1000km?

```
df.meanRadius
```

```
eName
Moon      33.0
Phobos    33.0
Deimos    33.0
Io        1821.5
Europa    1560.8
...
S/2017 J 8   0.5
S/2017 J 9   1.5
Ersa       1.5
Ultima Thule 33.0
101955 Bennu 33.0
Name: meanRadius, Length: 265, dtype: float64
```

- Extracting data

```
df.loc[df.meanRadius > 1000].meanRadius
```

```
eName
Io          1821.5000
Europa      1560.8000
Ganymede    2631.2000
Callisto    2410.3000
Tethys      1066.0000
```

- Masking data

```
df.where(df.meanRadius > 1000).meanRadius
```

```
eName
Moon        NaN
Phobos      NaN
Deimos      NaN
Io          1821.5
Europa      1560.8
```

Remove rows that contain nans

```
df.bondAlbido.dropna()
```

Filter data

Multiconditions:

AND: &

```
df.loc[(df.meanRadius > 10000) & (df.isPlanet)]
```

		isPlanet	semimajorAxis	perihelion	aphelion
eName					
Uranus	True	2870658186	2147483647	2147483647	
Neptune	True	4498396441	2147483647	2147483647	
Jupiter	True	778340821	740379835	816620000	
Saturn	True	1426666422	1349823615	1503509229	

OR: |

```
df.loc[(df.meanRadius > 10000) | (df.isPlanet)]
```

		isPlanet	semimajorAxis	perihelion	aphelion
eName					
1 Ceres	True	413690250	382620000	445410000	
136199 Eris	True	10180122852	2147483647	2147483647	
Uranus	True	2870658186	2147483647	2147483647	
Pluto	True	5906440628	2147483647	2147483647	

Numpy-like operations: how to

- Operation on one series:
`series.mean()`

```
df.meanRadius.mean()
```

```
3481.5921071698117
```

- Ignores nans by default `np.nanmean`
- Works on **Series** and **Dataframe**

```
df.max(numeric_only=True)
```

isPlanet	True
semimajorAxis	152000000000
perihelion	2147483647
aphelion	2147483647
eccentricity	0.7512
inclination	179.8

- Operation between aligned series:
`series1 + series2`

```
df.mass_kg / df.volume
```

Numpy-like operations: examples of existing operations

Statistical operations returning **one** metric

<code>mean</code>	Mean value	<code>min</code>	Lowest value	<code>max</code>	Highest value	<code>median</code>	Median value
<code>count</code>	Number of non nan values	<code>skew</code>	Skewness of the series	<code>kurt</code>	Kurtosis of the series	<code>mode</code>	Mode of the series
<code>var</code>	Variance of the series	<code>quantile</code>	Find quantiles	<code>sum</code>	Sum all elements in the series	<code>std</code>	Standard deviation

Some operations returning **multiple** values

<code>rank</code>	Ranks each elements in the series	<code>unique</code>	Returns unique values of elements	<code>cumsum</code>	Cumulative sum of the elemens	<code>cumprod</code>	Cumulative product of the elements
-------------------	-----------------------------------	---------------------	-----------------------------------	---------------------	-------------------------------	----------------------	------------------------------------

Get location of specific elements

<code>idxmin</code>	Get the index of the minimum value	<code>idxmax</code>	Get the index of the maximum value	<code>argmin</code>	Get the position of the minimum value	<code>argmax</code>	Get the position of the maximum value
---------------------	---	---------------------	---	---------------------	--	---------------------	--

Custom and multiple aggregations

- Apply **custom** function

```
def second_highest(data):
    sorted_data = data.sort_values(ascending=False)
    return sorted_data.iloc[1]
df.apply(second_highest).meanRadius
```

69911.0

- Apply **multiple** functions with **agg**

```
df.meanRadius.agg(['mean', 'std', 'var', 'max', 'min'])
```

```
mean      3.481592e+03
std       4.314101e+04
var        1.861146e+09
max       6.963420e+05
min       3.000000e-01
Name: meanRadius, dtype: float64
```

- Implemented **describe**

```
df.describe()
```

	semimajorAxis	perihelion	aphelion
count	2.650000e+02	2.650000e+02	2.650000e+02
mean	1.330740e+09	5.999915e+07	6.149610e+07
std	1.048573e+10	3.332062e+08	3.365186e+08
min	0.000000e+00	0.000000e+00	0.000000e+00
25%	2.946720e+05	0.000000e+00	0.000000e+00
50%	2.015529e+07	0.000000e+00	0.000000e+00
75%	2.392800e+07	0.000000e+00	0.000000e+00
max	1.520000e+11	2.147484e+09	2.147484e+09

Big warning: no loops !

- Many different available operations
- Potential to use **any** custom operation

If your code looks like this ... **DON'T!**

```
results = []
for i in range(df.index.size):
    line = df.iloc[i]
    results.append(line.meanRadius**2)
```

Do this instead !

```
results = df.meanRadius**2
```

1000x faster

**There is (almost) always a way to replace a
slow loop by a fast pandas function**

Sorting data

- Sort by index: `df.sort_index()`

- Sort by one of the columns

```
df.sort_values('meanRadius', ascending=False).meanRadius
```

eName	meanRadius
Sun	696342.0
Jupiter	69911.0
Saturn	58232.0
Uranus	25362.0
Neptune	24622.0

Grouping data using groupby

- We can create **groups** depending on a **column** and apply operations on each of the groups **separately**
- Using bins : pd.cut

```
bins = pd.cut(df.meanRadius, [0, 1000, 10000, 100000])
bins
```

eName	bins
Moon	(0, 1000]
Phobos	(0, 1000]
Deimos	(0, 1000]
Io	(1000, 10000]
Europa	(1000, 10000]

```
df.groupby('isPlanet').meanRadius.mean()
```

isPlanet	meanRadius
False	2867.689286
True	15381.862185

Name: meanRadius, dtype: float64

```
df.groupby(bins).gravity.mean()
```

meanRadius	gravity
(0, 1000]	0.013682
(1000, 10000]	2.128938
(10000, 100000]	13.812500

Name: gravity, dtype: float64

Grouping data using groupby

- **Multiple** criteria

```
df.groupby([df.isPlanet, bins]).meanRadius.max()

  isPlanet  meanRadius
  False      (0, 1000]      788.9000
              (1000, 10000]    2631.2000
              (10000, 100000]   NaN
  True       (0, 1000]      725.0000
              (1000, 10000]     6371.0084
              (10000, 100000]   69911.0000
Name: meanRadius, dtype: float64
```

Combining dataframes

- Concatenate using the indices

```
pd.concat([df1,df2])
```

df1

	method	number	orbital_period	mass	distance	year
0	Radial Velocity	1	269.300	7.10	77.40	2006
1	Radial Velocity	1	874.774	2.21	56.95	2008
2	Radial Velocity	1	763.000	2.60	19.84	2011

df2

	method	number	orbital_period	mass	distance	year
3	Radial Velocity	1	326.03	19.4	110.62	2007
4	Radial Velocity	1	516.22	10.5	119.47	2009
5	Radial Velocity	1	185.84	4.8	76.39	2008



	method	number	orbital_period	mass	distance	year
0	Radial Velocity	1	269.300	7.10	77.40	2006
1	Radial Velocity	1	874.774	2.21	56.95	2008
2	Radial Velocity	1	763.000	2.60	19.84	2011
3	Radial Velocity	1	326.030	19.40	110.62	2007
4	Radial Velocity	1	516.220	10.50	119.47	2009
5	Radial Velocity	1	185.840	4.80	76.39	2008

Combining dataframes

- Concatenate using columns

```
pd.concat([df1,df2], axis=1)
```

	method	number	orbital_period
0	Radial Velocity	1	269.300000
1	Radial Velocity	1	874.774000
2	Radial Velocity	1	763.000000
3	Radial Velocity	1	326.030000
4	Radial Velocity	1	516.220000

df1

	mass	distance	year
0	7.10	77.40	2006
1	2.21	56.95	2008
2	2.60	19.84	2011
3	19.40	110.62	2007
4	10.50	119.47	2009

df2



	method	number	orbital_period	mass	distance	year
0	Radial Velocity	1	269.300000	7.10	77.40	2006
1	Radial Velocity	1	874.774000	2.21	56.95	2008
2	Radial Velocity	1	763.000000	2.60	19.84	2011
3	Radial Velocity	1	326.030000	19.40	110.62	2007
4	Radial Velocity	1	516.220000	10.50	119.47	2009

One cool operation

- Compute correlation between all variables

df.corr(numeric_only=True)					
	isPlanet	semimajorAxis	perihelion	aphelion	eccentricity
isPlanet	1.000000	0.037617	0.792468	0.802667	
semimajorAxis	0.037617	1.000000	0.065412	0.064542	
perihelion	0.792468	0.065412	1.000000	0.999449	
aphelion	0.802667	0.064542	0.999449	1.000000	
eccentricity	-0.032432	-0.089975	0.012518	0.011152	

Part II summary: first operations

Some of the available methods:

Access data using index/columns + filters

`df.loc[index]` `df.loc[condition]`

Numpy-like **operations**

`df.mean()` `df.apply(function)`

Sort data by index/values

`df.agg(['std'])`

Operations on **groups** of data

`df.sort_values()` `df.sort_index()`

Combine multiple dataframes

`df.groupby(columns).mean()`

Correlation

`pd.concat([df1, df2])`

`df.corr()`

And so much more ...

Bonus: « one-liners »

- Each method usually returns a **new pandas object** (Series, DataFrame)
- We can keep **adding** more and more operations

```
df.loc[df.gravity > df.gravity.loc['Mercury']]\  
    .groupby(pd.cut(df.semimajorAxis_AU, [0,5,40]))\  
    .meanRadius\  
    .agg(['mean','std'])
```

	mean	std
semimajorAxis_AU		
(0, 5]	5270.769467	1637.026147
(5, 40]	44531.750000	23062.817411

Part II: Practicals

Go to the jupyter notebook



Part III: Plotting


Plotting interface, integration with matplotlib and seaborn

Plotting interface

- Direct integration of high level **plotting libraries** such as **matplotlib**


```
df_planet.plot.scatter(x='gravity', y='density', figsize=(3,3))
```

```
<AxesSubplot:xlabel='gravity', ylabel='density'>
```



Other backends possible (**Plotly**, **bokeh**, **altair**,...)

```
pd.options.plotting.backend = "plotly"
```




Plotting statistics for Series

- Integrated plotting methods:

```
series.plot.[type of plot]()
```


```
penguins.bill_length_mm.plot.hist()
```

```
<AxesSubplot:ylabel='Frequency'>
```




```
penguins.bill_length_mm.plot.kde()
```

```
<AxesSubplot:ylabel='Density'>
```



```
penguins.bill_length_mm.plot.box()
```

```
<AxesSubplot:>
```




- Pandas calls the **matplotlib** function: same arguments

Plotting categorical data with barplots


```
df_planet.density.sort_values().plot.bar(figsize=(4,2))
```

```
<AxesSubplot:xlabel='eName'>
```




```
df_planet[['aphelion','perihelion']].plot.barh(figsize=(3,3))
```

```
<AxesSubplot:ylabel='eName'>
```




Plotting relations between data: scatter/hexbin

```
df_planet.plot.scatter(x='semimajorAxis', y='density', figsize=(4,3))  
<AxesSubplot:xlabel='semimajorAxis', ylabel='density'>
```



```
penguins.plot.hexbin(x='bill_length_mm', y='body_mass_g',  
gridsize=20)
```

```
<AxesSubplot:xlabel='bill_length_mm', ylabel='body_mass_g'>
```



Plotting functions available

kind : str

The kind of plot to produce:

- 'line' : line plot (default)
- 'bar' : vertical bar plot
- 'barh' : horizontal bar plot
- 'hist' : histogram
- 'box' : boxplot
- 'kde' : Kernel Density Estimation plot
- 'density' : same as 'kde'
- 'area' : area plot
- 'pie' : pie plot
- 'scatter' : scatter plot (DataFrame only)
- 'hexbin' : hexbin plot (DataFrame only)

Integrating within a matplotlib workflow

Setup figure and axes:


```
import matplotlib.pyplot as plt  
fig, ax = plt.subplots(figsize=(4,4))
```

Plot with pandas on existing axis:

```
df_planet.plot.scatter(x='semimajorAxis',  
                      y='density',  
                      marker='x',  
                      ax=ax) ←
```

Change title, labels, limits, etc...

```
ax.set_title('Planets of the solar system')  
ax.set_xlim(0,6)  
ax.set_xlabel('Semi-major axis of the planet [km]e10')  
ax.set_ylabel('Density of the planet [1]')
```



What about seaborn?




seaborn

Visualization library
optimised for dataframes

Coming **soon** in the
Atelier numérique de l'OMP

```
import seaborn as sns
sns.relplot(data=df_planet_type,
             x='semimajorAxis', y='density',
             hue='planet_type', palette='Set2',
             size='mass_kg', sizes=(50,200),
             style=pd.Series(df_planet_type.index=='Earth', index=df_planet_type.index).rename('isEarth'),
             markers=['o','^'],
             )
```

<seaborn.axisgrid.FacetGrid at 0x21b03227c40>



Part III: Practicals

Go to the jupyter notebook



Part IV: Timeseries

Example of data

- Series or DataFrame with **datetime index**
- Constant time step or not
- *Example:* temperature logger in the ocean

Temperature
2019-04-13 14:00:00 28.4728
2019-04-13 14:00:10 28.4780
2019-04-13 14:00:20 28.4818
2019-04-13 14:00:30 28.5127
2019-04-13 14:00:40 28.5199
...
2020-10-02 03:19:09 28.6340
2020-10-02 03:19:19 28.6339
2020-10-02 03:19:29 28.6339
2020-10-02 03:19:39 28.6341
2020-10-02 03:19:49 28.6339

How to get a datetime index?

- Create a new index:

```
pd.date_range(start = '2014', end = '2018', freq = 'YS')
```

```
DatetimeIndex(['2014-01-01', '2015-01-01', '2016-01-01', '2017-01-01',
                 '2018-01-01'],
                dtype='datetime64[ns]', freq='YS-JAN')
```

```
pd.date_range(start = '2014-01-03 12:00:00', freq = 'min', periods = 10)
```

```
DatetimeIndex(['2014-01-03 12:00:00', '2014-01-03 12:01:00',
                 '2014-01-03 12:02:00', '2014-01-03 12:03:00',
                 '2014-01-03 12:04:00', '2014-01-03 12:05:00',
                 '2014-01-03 12:06:00', '2014-01-03 12:07:00',
                 '2014-01-03 12:08:00', '2014-01-03 12:09:00'],
                dtype='datetime64[ns]', freq='min')
```

- From strings

```
pd.to_datetime(['2010','2014','2023'])
```

```
pd.to_datetime(['2010-01-03','2014-02-12','2023-05-05'])
```


- Frequencies?

- YS, Q, MS, W, D, h, min, s, ms, us, ns
- Can also add a number: 10h

Plot the timeseries

```
df.Temperature.plot()
```

```
<AxesSubplot:>
```




Select data based on time

- Use `df.loc[time]` for **specific times**


```
df.Temperature.loc['2019-05'].plot()
```

```
<AxesSubplot:>
```




```
df.Temperature.loc['2019-05-11'].plot()
```

```
<AxesSubplot:>
```



```
df.Temperature.loc['2019-05-11 3H'].plot()
```

```
<AxesSubplot:>
```




Select data based on time

- Use `df.loc[time1:time2]` for **intervals**


```
df.Temperature.loc['2019-05':'2019-09'].plot()
```

```
<AxesSubplot:>
```



```
df.Temperature.loc['2019-05-11':'2019-05-15'].plot()
```

```
<AxesSubplot:>
```



Resampling

- Use the same **frequencies** as in date_range:

```
| df.Temperature.resample('H')
```


- Returns an object that **expects an operation**
 - mean, max, quantile, count, ... or apply()

```
| df.Temperature.resample('D').mean()
```

2019-04-13	29.869072
2019-04-14	29.883743
2019-04-15	29.671858
2019-04-16	29.371685
2019-04-17	29.896633

```
df.Temperature.plot()  
df.Temperature.resample('D').mean().plot(lw=3)
```

```
<AxesSubplot:>
```




Rolling windows

- Keeps the **same sampling rate**
- Returns an object as in groupby: **expects an operation**
 - `mean`, `max`, `quantile`, `count`, ... or `apply()`


```
s.plot()  
s.rolling(50, center=True).max().plot(figsize=(4,3))
```

<AxesSubplot:>



```
s.plot()  
s.rolling('2D').mean().plot(figsize=(4,3))
```

<AxesSubplot:>



Dealing with missing data

- Can we get **daily** data?


Temperature	
2019-04-14 01:04:09	29.8508
2019-04-14 16:16:29	29.8557
2019-04-15 05:10:19	30.0335
2019-04-15 07:51:29	29.4874
2019-04-15 14:08:49	29.8838



.resample('D').mean()


```
df_irregulier.Temperature.resample('D').mean().plot()
```

```
<AxesSubplot:>
```



```
df_irregulier.Temperature.resample('D').mean().interpolate().plot()
```

```
<AxesSubplot:>
```



.interpolate()



Access date information


- Many attributes available for a **DatetimeIndex**
 - month, year, dayofweek, is_month_start, is_leap_year, ...

```
df.index.month|
```

```
Int64Index([13, 13, 13, 13, 13, 13, 13, 13, 13, 13,
            ...
            30, 30, 30, 30, 30, 30, 30, 30, 30, 30],
            dtype='int64', length=677521)
```

Groupby for temporal data?

- Daily temperature cycle:



Part IV summary

series/dataframe with datetime index (numpy datetime/pandas date_range, ...)

Some available methods:

Temporal selection: `df.loc["2012-10-04"]`

Period selection: `df.loc["2008-07" : "2015-07"]`

Resampling: `df.resample("D").mean()`

Rolling windows: `df.rolling(5).mean()` `df.rolling('2D').mean()`

Fill the nans: `df.interpolate()`

Datetime infos: `df.index.hour` `df.index.day_of_year`

Part IV: Practicals

Go to the jupyter notebook

