

xarray

# Introduction to xarray

*Managing and analyzing multidimensional datasets*

# Programme: four blocks

Introduction to  
DataArrays and  
Datasets,  
reading/writing

Plotting capacities

First analyses

**Scaling analysis  
with dask**



```
import xarray as xr
```

# Part I: xarray objects

DataArrays and Datasets

# Comparison with pandas

Multivariable objects, aligned on similar axes/indices

- **pd.DataFrame**
- **xr.Dataset**

Single variable object:

- **pd.Series**
- **xr.DataArray**

Single variable object :


- **pd.Series**
- **xr.DataArray**

Single variable object :

- **pd.Series**
- **xr.DataArray**

# xr.DataArray: presentation

- Array of values representing a **unique** variable:
  - « Wrapping » around a numpy array



# xr.DataArray: presentation

- Exemple of a `DataArray`

```
xarray.DataArray 'thetao' (time: 312, depth: 13, latitude: 157, longitude: 265)
```

[168749880 values with dtype=float32]

▼ Coordinates:

<b>depth</b>	(depth)	float32	13.47 15.81 18.5 ... 77.85 92.33	 
<b>latitude</b>	(latitude)	float32	-5.0 -4.917 -4.833 ... 7.917 8.0	 
<b>time</b>	(time)	datetime64[ns]	1993-01-16T12:00:00 ... 2018-12-	 
<b>longitude</b>	(longitude)	float32	-180.0 -179.9 ... -158.1 -158.0	 

▼ Attributes:

long_name :	Temperature
standard_name :	sea_water_potential_temperature
units :	degrees_C
unit_long :	Degrees Celsius
cell_methods :	area: mean
_ChunkSizes :	[ 1 7 341 720]

# xr.DataArray: creating an object

- Source: an n-dimensional **numpy array** + n lists of **coordinates**

```
data = np.array([[1, 2, 3],  
                [5, 6, 7]])  
longitude = [10, 15, 20]  
latitude = [0, 5]
```

- Use **xr.DataArray()**

```
dataArray = xr.DataArray(data,  
                         dims = ['latitude', 'longitude'],  
                         coords = {'longitude':longitude,  
                                    'latitude':latitude  
                         })
```

```
xarray.DataArray (latitude: 2, longitude: 3)  
  
array([[1, 2, 3],  
       [5, 6, 7]])  
▼ Coordinates:  
  longitude      (longitude) int32 10 15 20  
  latitude       (latitude)  int32 0 5  
► Attributes: (0)
```

# xr.DataArray: adding attributes

- `DataArray.attrs` is a **dictionary** that can be modified

```
dataarray.attrs['units'] = '°C'  
dataarray.attrs['description'] = "Température de l'OMP"
```

```
xarray.DataArray (latitude: 2, longitude: 3)  
  
  array([[1, 2, 3],  
         [5, 6, 7]])  
  ▼ Coordinates:  
    longitude      (longitude) int32 10 15 20  
    latitude       (latitude)   int32 0 5  
  ▼ Attributes:  
    units :          °C  
    description :  Température de l'OMP
```

# xr.DataArray: accessing the data

## Xarray basics ! Access data using coordinates

```
dataarray.sel(longitude = 15)
```

```
xarray.DataArray  (latitude: 2)
```

```
  array([2, 6])
```

▼ Coordinates:

longitude	( )	int32	15
-----------	-----	-------	----

latitude	(latitude)	int32	0 5
----------	------------	-------	-----

# xr.DataArray: accessing the data

- Also works for a list of **multiple** coordinates

```
dataarray.sel(longitude = [15, 10], latitude = 5)
```

- Select using **position** (like in numpy)


```
dataarray.isel(longitude = [0,2])
```

- Access the numpy **array**

```
dataarray.values
```

# xr.Dataset: presentation

- Multiple DataArrays with shared dimensions



# xr.Dataset: creating an object

- Sources:
  - An ensemble of DataArrays
- `xr.Dataset()`

```
dataset = xr.Dataset({"temperature":data_temperature,  
                      "salinity":data_salinity,  
                      "precipitation":data_precipitation})
```

# xr.Dataset: access the data

- Access one DataArray

```
dataset.temperature  
dataset['temperature']
```

- Other access are **similar** to a DataArray:
  - `dataset.sel(latitude=...)`
  - `dataset.isel(longitude=[..])`

# xr.Dataset: append data

- Add a DataArray as a **new variable**
  - Add a variable as if you wanted to access it:

```
dataset['ensoleillement'] = dataarray
```

- This can also be used to **replace** the values of a given DataArray

# Operations on coordinates

- Simple operations between DataArrays based on **coordinates**:

```
DataArray1**2 + DataArray2/12
```

- Create **new variables** in a dataset from **existing variables**:

```
ds['new_variable'] = ds['old_v1']**2 + ds['old_v2']/12
```

# Combine dataarrays or datasets

- Concatenate along **existing dimensions**
  - Example: different times but similar geographical grids

```
xr.concat([ds1, ds2], dim = 'latitude')
```


- Concatenate along **new dimensions**

```
xr.concat([ds1, ds2, ds3, ds4], dim = 'modele')
```

# Combine dataarrays or datasets

ds1


xarray.DataArray 'thetao' (**latitude**: 157, **longitude**: 265)



```
xr.concat([ds1, ds2, ds3], dim='time')
```

xarray.DataArray 'thetao' (time: 3, **latitude**: 157, **longitude**: 265)

`xr.concat([ds1,ds2,ds3],  
dim='time')`



# Reading/writing files: supported formats

- Xarray supports many different format for reading and writing
  - [netCDF](#)
  - Other examples
    - GRIB
    - geoTIFF
    - zarr
    - ...
- Optimized for [gridded data](#)
- Can open directly [online](#) (OpenDAP, cloud buckets, ...)

# Reading/writing files: reading a file

- To **open** a netcdf file, use `xr.open_dataset()`

```
ds = xr.open_dataset('../data/GLORYS_ocean-temp-currents_1993-2019.nc')
```

- There is also `open_dataarray` if there is only **one** variable
- **Reading multiple files with `xr.open_mfdataset`:**

- With a pattern

```
dataset = xr.open_mfdataset("../data/GLORYS*.nc")
```

- With a list of paths

```
dataset = xr.open_mfdataset(["/path/to/file1.nc", "/path/to/file2.nc"])
```

# Reading/writing files: writing a file

- To **save** a file, use `Dataset.to_format()`

```
ds.to_netcdf('path_to_file.nc')
```

# Part I: Summary of xarray objects

Wrap numpy arrays with:

- **Dimensions** that have names (lon, lat, time, depth, altitude, ...)
- **Coordinates** that have values
- **Attributes**

**DataArray**: for a unique variable (equivalent to pd.Series)

**Dataset**: for multiple variables on the same « grid » (equivalent to pd.DataFrame)

Access data using variable **names** and **coordinates**

```
ds.temperature.sel(lon=54, lat=[12,13,14])
```

# Part I: Practicals

Go to the jupyter notebook




# Part II: plotting capacities

# Plotting DataArrays

- As pandas, xarray integrates **matplotlib** (+ other backends)

```
ds.thetao.isel(time=0, depth=0).plot()  
<matplotlib.collections.QuadMesh at 0x183d4e58430>
```




- Plot shows **coordinates**, **dimensions** and **attributes**

# Plotting DataArrays: default behaviour

- Line plot for 1d data

```
ds.thetao.mean(['longitude','latitude','depth']).plot()
```


```
[<matplotlib.lines.Line2D at 0x183d4da0fa0>]
```



- Quadmesh for 2D data


```
ds.thetao.isel(time=0, depth=0).plot()
```

```
[<matplotlib.collections.QuadMesh at 0x183d4e58430>]
```




- Histogram if more than 2D

```
ds.thetao.plot();
```




# Plotting DataArrays: possible 2D plots


DataArray.plot.quadmesh()



DataArray.plot.contourf()



DataArray.plot.contour()



## Important key words:


<code>x='longitude'</code>	Variable in x	<code>levels=10</code>	Number of levels or list
<code>y='depth'</code>	Variable in y	<code>cmap='magma'</code>	Changes colormap
<code>yincrease=True</code>	Reverses y values	<code>robust=True</code>	Truncates extreme values

# Plotting DataArrays: reduce plot dimension

- Plot 2D data as **multiple line plots** with the keyword hue


```
ds.isel(longitude=0, latitude=0, depth=[0,5,10]).thetao.plot()
```

```
<matplotlib.collections.QuadMesh at 0x1d29d264d30>
```



```
ds.isel(longitude=0, latitude=0, depth=[0,5,10]).thetao.plot(hue='depth')
```


```
[<matplotlib.lines.Line2D at 0x1d29d31cf0>,
 <matplotlib.lines.Line2D at 0x1d29d31cd90>,
 <matplotlib.lines.Line2D at 0x1d29d2cc670>]
```



# Plotting DataArrays: reduce plot dimension

- Plot 3D data as **multiple 2D plots** with the keywords `col/row`

```
ds.mean('depth').isel(time=[0,12,24]).thetao.plot(col='time')
```




- Use `col_wrap=5` if there are **two many columns**

# Plotting datasets

You can plot relationship between **different variables** of the **same dataset**


`Dataset.plot.scatter(x,y)`

```
ds.isel(time=0, depth=0).plot.scatter(x='uo',
                                      y='vo',
                                      hue='thetao');
```




`Dataset.plot.quiver(x,y,u,v)`

```
ds.isel(time=0, depth=0).plot.quiver(x='longitude',
                                      y='latitude',
                                      u='uo',
                                      v='vo')
```



`Dataset.plot.streamplot(x,y,u,v)`

```
ds.isel(time=0, depth=0).plot.streamplot(x='longitude',
                                         y='latitude',
                                         u='uo',
                                         v='vo')
```



# Bonus: integration with cartopy

Import cartopy

```
import cartopy.crs as ccrs
```

Prepare figure

```
fig, ax = plt.subplots(  
    figsize=(6,2),  
    subplot_kw = {'projection':ccrs.Robinson(200)})
```

**Declare projection**


Plot data

```
data = ds.thetao.isel(time=0, depth=4)  
  
data.plot(ax=ax,  
          transform=ccrs.PlateCarree())
```

**Declare coordinate system**

Change limits/add coastline

```
ax.set_extent((155,285,-10,10))  
ax.coastlines()
```



# Part II: Practicals

Go to the jupyter notebook



# Part III: First analyses

- Data selection
- Statistical operations
- Aggregations, ...

# Part III: First analyses

- **Data selection**
- Statistical operations
- Aggregations, ...

# Selecting data

- Select by **coordinates**

```
ds.sel(latitude=-4.75)
```

- Select by **position**

```
ds.isel(depth=5)
```

- Unique selection

```
ds.sel(latitude=-4.75)
```

- Selection from list

```
ds.isel(depth=[5,8])      ds.isel(longitude = np.arange(3,9))
```

- Multicriteria

```
ds.isel(depth=[5,8], longitude = 12)
```

- Nearest neighbour

```
ds.sel(latitude=50, method='nearest')
```

# Selecting data

- **Range** of coordinates: `slice(beginning, end)`
  - Select **all values** between beginning and end

```
ds.sel(latitude=slice(-3,3))
```

- **Temporal** selection (similar to pandas):

- **Exact** selection

```
ds.sel(time='1993-01-16')
```

- Selection of a **period**

```
ds.sel(time='2015')
```

- Selection of a **time span**

---

```
ds.sel(time=slice('2001', '2015-04'))
```

# Selecting data

- **Masking** data with `Dataset.where`

- Selects the data that fulfill a certain **condition**, the rest will be NaN.
- The condition is **a boolean dataarray** (True/False) on the same coordinates

```
ds.where(ds.uo>0.1)    ds.where(ds.uo>0.1, other=999)
```

- When the condition is on the **coordinates**, **extract** the data with `drop=True`

```
ds.where(ds.latitude<-2, drop=True)
```

# Part III: First analyses

- Data selection
- **Statistical operations**
- Aggregations, ...

# Numpy-like operations

- Use `DataArray.operation('dimension')` for one object

```
ds.thetao.mean('depth')
```

```
| ds.thetao.mean(['longitude', 'latitude'])
```

- Same for Datasets    `ds.mean(['longitude', 'latitude'])`

- Generally the same operations as in pandas

- `mean()`, `sum()`, `min()`, `max()`, `median()`
- `idxmax()`, `idxmin()`, `argmin()`, `argmax()`
- `quantile([q1, q2, ...])`
- `count()`

# Some advanced methods

- **Difference** from one step to another

```
ds.diff('time')
```

- **Cumulative sum**

```
ds.cumsum('time')
```

- **Gradient, integral**

```
ds.differentiate('time')  
ds.integrate('time')
```

# Weighted operations

- Weighted operations are possible !
- The weights are a `DataArray` with similar dimensions/coordinates
  - Example: Global average weighted by cell area of a climate model

```
ds.weighted(cell_size).mean(['longitude','latitude'])
```

# Fitting a DataArray or a dataset

- Polynomial fits or general **curve fits** as with scipy/numpy

```
fit = ds.isel(depth=0).polyfit("time", deg=1)
```

xarray.Dataset

► Dimensions: (degree: 2, latitude: 157, longitude: 265)

▼ Coordinates:

degree	(degree)	int32 1 0
latitude	(latitude)	float64 -5.0 -4.917 -4.833 ... 7.917 8.0
longitude	(longitude)	float64 -180.0 -179.9 ... -158.1 -158.0

▼ Data variables:


vo_polyfit_coeffi...	(degree, latitude, longitude)	float64 -3.185e-20 -3.72e-20 ... 0.05072
thetao_polyfit_c...	(degree, latitude, longitude)	float64 1.37e-19 1.262e-19 ... 27.78 27.78
uo_polyfit_coeffi...	(degree, latitude, longitude)	float64 -5.6e-20 -5.629e-20 ... 0.105 0.106

► Attributes: (17)

```
def f(x, a,b):
    return a*x+b
fit = ds.isel(depth=0, longitude=0).curvefit("time",f)
```

```
fit.thetao_polyfit_coefficients.sel(degree=1).plot()
```

```
<matplotlib.collections.QuadMesh at 0x1d29a8f54c0>
```



# Big warning: no loops !

**There is (almost) always a way to replace a  
slow loop by a fast xarray function**

Example : create **new dimension** instead of loop

# BONUS: Applying custom function to some dimension

Define **custom function**:

```
def get_second_highest(data):
    sorted_data = np.sort(data)
    return sorted_data[-2]
```

Apply it to certain dimensions with `xr.apply_ufunc`:

```
xr.apply_ufunc(get_second_highest,
                 ds.uo,
                 input_core_dims=[['time']],
                 output_core_dims=[[]],
                 vectorize=True
                )
```

# Part III: First analyses

- Data selection
- Statistical operations
- **Aggregations, ...**

# Resampling and rolling windows: in time


In time: similar to pandas

Resampling:

```
ds.resample(time="2H").mean()
```

Rolling windows:

```
ds.rolling(time=120).median()
```




# Resampling and rolling windows: in other dimensions


**Resampling:** `ds.coarsen(longitude=20, latitude=20, boundary='pad').mean()`

**Rolling windows:** `ds.rolling(longitude=20, latitude=20, center=True).max()`


Raw DataArray



DataArray.rolling(lon=20, lat=20)




DataArray.coarsen(lon=20, lat=20)



# Bonus: regridding with xesmf

## xESMF: Universal Regridder for Geospatial Data

xESMF is a Python package for [regridding](#). It is



- **Curvilinear grid** (e.g. NEMO, global ocean model)
- **Rectilinear grid** (regular lon/lat grid)

# Interpolation

- Interpolate on **new coordinate values**

```
ds.interp(latitude=np.arange(-5, 5, 0.1), method='linear')
```

- Fill **missing values**:

```
ds.interpolate_na('latitude', method='cubic')
```

# Operations on groups

- Use `Dataset.groupby()` to compute on **separate groups** of data (as in pandas)
  - Apply operation on each group

```
small_ds.groupby(small_ds.uo//0.1).mean()
```

- **Temporal** groupby using `Dataset.groupby('time.XXX')`:

```
ds.groupby('time.month')
```

# Bonus: faster groupby with flox

## flox: Faster GroupBy reductions with Xarray

Tuesday, July 18th, 2023 (10 months ago)



Deepak Cherian

- Better and faster algorithms
- Optimized for parallel computing
- Works with dask (see end of the day...)

# Summary: first analyses

Select data with [coordinates](#)

```
ds.sel(longitude=12, latitude=slice(0,40), time="2012")
```

Mask data with [conditions](#)

```
ds.where(ds.temperature > 18)
```

Operation along [dimension](#)

```
ds.max("depth")
```

[Weighted operations](#)

```
ds.weighted(cell_size).mean(["longitude", "latitude"])
```

[Resampling](#) in time

```
ds.resample(time="2H").min()
```

Other [resampling](#)

```
ds.coarsen(longitude=10, latitude=5, boundary="trim").max()
```

[Rolling windows](#)

```
ds.rolling(depth=5, center=True)
```

[Interpolating](#) on new coordinates

```
ds.interp(longitude = [10, 20, 30], latitude=18)
```

# One-liners

- As in **pandas**, methods **return** an xarray object (**Dataset**, **DataArray**)
- Methods can be **chained** into one line:

```
ds.thetao.where(ds.uo>0.1)\\
    .resample(time='Y').mean()\\
    .sel(latitude=slice(-2,2))\\
    .mean('longitude')\\
    .integrate('depth')
```

- **ATTENTION:** it can be better to split for development...

# Easy conversion to pandas objects

- xarray objects can be **directly converted** to **pandas object**

```
ds.thetao.to_series()
```

```
ds.to_dataframe()
```

# Part III: Practicals

Go to the jupyter notebook



# Part IV: Scaling with dask

Handling out of memory datasets & parallel computing

# Scaling issues

```
xarray.DataArray 'thetao' (model: 100, time: 312, depth: 13, latitude: 157, longitude: 265)
```


- How to handle **large amount** of data (Here 63 GB)?
  - Need **large computers** ?
- Limited by **memory size** : can't load **more** than memory size

# Introducing chunks

- Split data into **unit boxes**

```
xarray.DataArray 'thetao' (model: 100, time: 312, depth: 13, latitude: 157, longitude: 265)
```

	Array	Chunk			
<b>Bytes</b>	62.86 GiB	154.72 MiB			
<b>Shape</b>	(100, 312, 13, 157, 265)	(100, 312, 13, 10, 10)			
<b>Count</b>	865 Tasks	432 Chunks	100	312	157
<b>Type</b>	float32	numpy.ndarray			




# Lazy computing

- When opening a file, only **metadata**, **coordinates** and **dimensions** are loaded
- Compute stuff ...
- The data are loaded into memory only when they are **really needed** (plotting, saving results, printing values)

# Parallel computing

- Operations on individual chunks
  - Sort operations
  - Computing them individually



# Quick setup

Providing **resource** information:

```
from dask.distributed import Client  
  
client = Client(n_workers=6)
```

Specify explicit **chunks**:

```
ds = xr.open_dataset("../data/GLORYS_ocean-temp-currents_1993-2019.nc",  
                    chunks = {'longitude':10, 'latitude':10})
```

Open **multiple files**: one chunk per file

```
ds = xr.open_mfdataset("../data/GLORYS_*.nc")
```

# Summary: why use dask?

- **Transparently** integrated with xarray
- Prevents **memory** issues
- Speeds up operations with **parallel computing**
- **Optimized for HPC and cloud**

# Live example ...